



# Database Systems

## Lecture 10: Object-Based Databases

**Dr. Momtazi**  
**[momtazi@aut.ac.ir](mailto:momtazi@aut.ac.ir)**

based on the slides of the course book



# Object-Based Databases

- **Complex Data Types and Object Orientation**
- Structured Data Types and Inheritance in SQL
- Array and Multiset Types in SQL
- Persistent Programming Languages
- Comparison of Object-Oriented and Object-Relational Databases



# Object-Relational Data Models

- Extend the relational data model by including object orientation and constructs to deal with added data types.
- Allow attributes of tuples to have complex types, including non-atomic values such as nested relations.
- Preserve relational foundations, in particular the declarative access to data, while extending modeling power.
- Upward compatibility with existing relational languages.



# Complex Data Types

- Motivation:
  - Permit non-atomic domains
  - Example of non-atomic domain: set of integers, or set of tuples
  - Allows more intuitive modeling for applications with complex data
  
- Intuitive definition:
  - allow relations whenever we allow atomic (scalar) values — relations within relations
  - Retains mathematical foundation of relational model
  - Violates first normal form.



# Example of a Nested Relation

- Example: library information system
- Each book has
  - title,
  - a list (array) of authors,
  - Publisher, with subfields *name* and *branch*, and
  - a set of keywords
- Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		( <i>name, branch</i> )	
Compilers	[Smith, Jones]	(McGraw-Hill, NewYork)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}



# 4NF Decomposition of Nested Relation

- Suppose for simplicity that title uniquely identifies a book
  - In real world ISBN is a unique identifier
- Decompose *books* into 4NF using the schemas:
  - $(title, author, position)$
  - $(title, keyword)$
  - $(title, pub\_name, pub\_branch)$
- 4NF design requires users to include joins in their queries.

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

*authors*

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

*keywords*

<i>title</i>	<i>pub_name</i>	<i>pub_branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

*books4*



# Complex Types and SQL

- Extensions introduced in SQL:1999 to support complex types:
  - Collection and large object types
    - ▶ Nested relations are an example of collection types
  - Structured types
    - ▶ Nested record structures like composite attributes
  - Inheritance
  - Object orientation
    - ▶ Including object identifiers and references
- Not fully implemented in any database system currently
  - But some features are present in each of the major commercial database systems
    - ▶ Read the manual of your database system to see what it supports



# Object-Based Databases

- Complex Data Types and Object Orientation
- **Structured Data Types and Inheritance in SQL**
- Array and Multiset Types in SQL
- Persistent Programming Languages
- Comparison of Object-Oriented and Object-Relational Databases





# Structured Types and Inheritance in SQL

- **Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

```
create type Name as  
  (firstname      varchar(20),  
   lastname      varchar(20))  
final
```

```
create type Address as  
  (street        varchar(20),  
   city          varchar(20),  
   zipcode       varchar(20))  
not final
```

- Note: **final** and **not final** indicate whether subtypes can be created



# Structured Types and Inheritance in SQL

- Structured types can be used to create tables with composite attributes

```
create table person (  
    name      Name,  
    address Address,  
    dateOfBirth date)
```

- Dot notation used to reference components: *name.firstname*



# Structured Types (cont.)

- **User-defined row types**

```
create type PersonType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

- Can then create a table whose rows are a user-defined type  
**create table** *person* **of** *PersonType*

- Alternative using **unnamed row types**.

```
create table person_r(  
    name    row(firstname varchar(20),  
                lastname varchar(20)),  
    address row(street    varchar(20),  
                city      varchar(20),  
                zipcode varchar(20)),  
    dateOfBirth date)
```



# Methods

- Can add a method declaration with a structured type.

**method** *ageOnDate* (*onDate* **date**)

**returns** interval year

- Method body is given separately.

**create instance method** *ageOnDate* (*onDate* **date**)

**returns** interval year

**for** *PersonType*

**begin**

**return** *onDate* - **self.dateOfBirth**;

**end**

- We can now find the age of each person:

**select** *name.lastname*, *ageOnDate* (**current\_date**)

**from** *person*



# Constructor Functions

- **Constructor functions** are used to create values of structured types
- E.g.  
**create function** *Name*(*firstname* **varchar**(20), *lastname* **varchar**(20))  
**returns** *Name*  
**begin**  
    **set** **self**.*firstname* = *firstname*;  
    **set** **self**.*lastname* = *lastname*;  
**end**
- To create a value of type *Name*, we use  
**new** *Name*('John', 'Smith')
- Normally used in insert statements  
**insert into** *Person* **values**  
    (**new** *Name*('John', 'Smith'),  
    **new** *Address*('20 Main St', 'New York', '11001'),  
    **date** '1960-8-22');



# Type Inheritance

- Suppose that we have the following type definition for people:

```
create type Person  
  (name varchar(20),  
   address varchar(20))
```

- Using inheritance to define the student and teacher types

```
create type Student  
under Person  
  (degree varchar(20),  
   department varchar(20))
```

```
create type Teacher  
under Person  
  (salary integer,  
   department varchar(20))
```

- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration



# Multiple Type Inheritance

- SQL:1999 and SQL:2003 do not support multiple inheritance
- If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:

```
create type Teaching Assistant  
under Student, Teacher
```

- To avoid a conflict between the two occurrences of *department* we can rename them

```
create type Teaching Assistant  
under  
  Student with (department as student_dept),  
  Teacher with (department as teacher_dept)
```

- Each value must have a **most-specific type**



# Table Inheritance

- Tables created from subtypes can further be specified as **subtables**
- E.g.
  - create table** *people* **of** *Person*;
  - create table** *students* **of** *Student* **under** *people*;
  - create table** *teachers* **of** *Teacher* **under** *people*;
- Tuples added to a subtable are automatically visible to queries on the supertable
  - E.g. query on *people* also sees *students* and *teachers*.
  - Similarly updates/deletes on *people* also result in updates/deletes on subtables
  - To override this behaviour, use “**only people**” in query





# Object-Based Databases

- Complex Data Types and Object Orientation
- Structured Data Types and Inheritance in SQL
- **Array and Multiset Types in SQL**
- Persistent Programming Languages
- Comparison of Object-Oriented and Object-Relational Databases



# Array and Multiset Types in SQL

- SQL supports two collection types:
  - Arrays
    - ▶ array types were added in SQL:1999
  - Multisets
    - ▶ multiset types were added in SQL:2003 .
- Multiset is an unordered collection, where an element may occur multiple times.
- Multisets are like sets, except that a set allows each element to occur at most once.
- Unlike elements in a multiset, the elements of an array are ordered
  - So we can distinguish the first item from the second item, and so on.



# Array and Multiset Types in SQL

- Example of array and multiset declaration:

```
create type Publisher as  
  (name          varchar(20),  
   branch       varchar(20));  
  
create type Book as  
  (title          varchar(20),  
   author_array  varchar(20) array [10],  
   pub_date      date,  
   publisher     Publisher,  
   keyword-set   varchar(20) multiset);  
  
create table books of Book;
```



# Creation of Collection Values

- Array construction

```
array ['Silberschatz', 'Korth', 'Sudarshan']
```

- Multisets

```
multiset ['computer', 'database', 'SQL']
```

- To create a tuple of the type defined by the books relation:

```
('Compilers', array['Smith', 'Jones'],  
  new Publisher ('McGraw-Hill', 'New York'),  
  multiset ['parsing', 'analysis'] )
```

- To insert the preceding tuple into the relation books

```
insert into books  
values
```

```
('Compilers', array['Smith', 'Jones'],  
  new Publisher ('McGraw-Hill', 'New York'),  
  multiset ['parsing', 'analysis'] );
```



# Querying Collection-Valued Attributes

- To find all books that have the word “database” as a keyword,

```
select title  
from books  
where 'database' in (unnest(keyword-set))
```

- We can access individual elements of an array by using indices
  - E.g.: If we know that a particular book has three authors, we could write:

```
select author_array[1], author_array[2], author_array[3]  
from books  
where title = 'Database System Concepts'
```

- To get a relation containing pairs of the form “title, author\_name” for each book and each author of the book

```
select B.title, A.author  
from books as B,  
      unnest (B.author_array) as A (author)
```



# Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.
- E.g.  
**select** *title*, *A* **as** *author*, *publisher.name* **as** *pub\_name*,  
          *publisher.branch* **as** *pub\_branch*, *K.keyword*  
**from** *books* **as** *B*,  
  
          **unnest**(*B.author\_array* ) **as** *A* (*author* ),  
          **unnest** (*B.keyword\_set* ) **as** *K* (*keyword* )
- Result relation *flat\_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



# Nesting

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset
- To nest the *flat\_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as publisher,  
        collect (keyword) as keyword_set  
from flat_books  
groupby title, author, publisher
```

<i>title</i>	<i>author</i>	<i>publisher</i>	<i>keyword_set</i>
		(pub_name, pub_branch)	
Compilers	Smith	(McGraw-Hill, New York)	{parsing, analysis}
Compilers	Jones	(McGraw-Hill, New York)	{parsing, analysis}
Networks	Jones	(Oxford, London)	{Internet, Web}
Networks	Frick	(Oxford, London)	{Internet, Web}



# Nesting

- To nest on both authors and keywords:

```
select title,  
        collect (author) as author_set,  
        Publisher (pub_name, pub_branch) as publisher,  
        collect (keyword) as keyword_set  
from   flat_books  
group by title, publisher
```





# Nesting

- To nest on both authors and keywords using set():

```
select title,  
        set (author) as author_set,  
        Publisher (pub_name, pub_branch) as publisher,  
        set (keyword) as keyword_set  
from   flat_books  
group by title, publisher
```



# Nesting

- To nest on both authors and keywords using subqueries:

```
select title,  
      (select author  
        from flat-books as M  
        where M.title=O.title) as author-set,  
      Publisher (pub-name, pub-branch) as publisher,  
      (select keyword  
        from flat-books as N  
        where N.title = O.title ) as keyword-set  
from flat-books as O
```



# Nesting

- To nest on both authors and keywords using subqueries:

```
select title,  
      array( select author  
            from authors as A  
            where A.title = B.title  
            order by A.position) as author_array,  
      Publisher(pub_name, pub_branch) as publisher,  
      multiset( select keyword  
                from keywords as K  
                where K.title = B.title) as keyword_set,  
from books4 as B;
```



# Object-Based Databases

- Complex Data Types and Object Orientation
- Structured Data Types and Inheritance in SQL
- Array and Multiset Types in SQL
- **Persistent Programming Languages**
- Comparison of Object-Oriented and Object-Relational Databases



# Persistence of Objects

- Object-oriented programming languages already have a concept of objects
  - A type system to define object types
  - Constructs to create objects.
- These objects are transient
  - They vanish when the program terminates
  - Similar to variables in a Java or C program vanish when the program terminates.
- To turn such a language into a database programming language, the first step is to provide a way to make objects persistent.



# Persistence of Objects

- Several approaches have been proposed:
  - **Persistence by class** - explicit declaration of persistence
  - **Persistence by creation** - special syntax to create persistent objects
  - **Persistence by marking** - make objects persistent after creation
  - **Persistence by reachability** - object is persistent if it is declared explicitly to be so or is reachable from a persistent object



# Object-Based Databases

- Complex Data Types and Object Orientation
- Structured Data Types and Inheritance in SQL
- Array and Multiset Types in SQL
- Persistent Programming Languages
- **Comparison of Object-Oriented and Object-Relational Databases**



# Object-Relational Mapping

- **Object-Relational Mapping (ORM)** systems built on top of traditional relational databases
- Implementer provides a mapping from objects to relations
- Objects can be retrieved from database
  - System uses mapping to fetch relevant data from relations and construct objects
  - Updated objects are stored back in database by generating corresponding update/insert/delete statements





# Comparison of O-O and O-R Databases

- Object-relational databases are object-oriented data-bases built on top of the relation model
- Object-oriented databases are built around persistent programming languages
- Object-relational mapping systems build an object layer on top of a traditional relational database
- Each of these approaches targets a different market.



# Comparison of O-O and O-R Databases

- Object-relational systems aim at making data modeling and querying easier by using complex data types.
  - Typical applications include storage and querying of complex data, including multimedia data.
  - SQL , however, imposes a significant performance penalty for certain kinds of applications that run primarily in main memory, and that perform a large number of accesses to the database.



# Comparison of O-O and O-R Databases

- Persistent programming languages target such applications that have high performance requirements.
  - They provide low-overhead access to persistent data and eliminate the need for data translation if the data are to be manipulated by a programming language.
  - However, they are more susceptible to data corruption by programming errors, and they usually do not have a powerful querying capability.
  - Typical applications include CAD databases.



# Comparison of O-O and O-R Databases

- Object-relational mapping systems allow programmers to build applications using an object model, while using a traditional database system to store the data.
  - They combine the robustness of widely used relational database systems, with the power of object models for writing applications.
  - However, they suffer from overheads of data conversion between the object model and the relational model used to store data.



# Comparison of O-O and O-R Databases

## ■ Relational systems

- simple data types, powerful query languages, high protection.

## ■ Object-relational systems

- complex data types, powerful query languages, high protection.

## ■ Persistent-programming-language-based OODBs

- complex data types, integration with programming language, high performance.

## ■ Object-relational mapping systems

- complex data types integrated with programming language, but built as a layer on top of a relational database system



**Questions?**